

Project 3: Working with MapReduce

DS 730

Overview

In this project, you will be working with input, output, Python and the MapReduce framework. You will be writing one mapper and one reducer per problem to solve a few different problems. None of the code you write will be multithreaded but the code that you create for this project will eventually be plugged into Hadoop. Hadoop will automatically parallelize your program for you. In short, you will be writing rather simple code that works on 1 machine right now. Eventually we will be using tools to help us parallelize the code (without having to change anything) so that we can solve our problems much faster.

As said in the notes, the map function is **stateless**. The intermediate output the mapper produces must only depend on the current input value. Said differently, if you are storing your rows for future lookup in your map function, you are doing something wrong. This is a very subtle but very important concept. If you are unsure if you are doing your mapping correctly, ask me.

Similarly, the output the reducer produces must only depend on a particular key. There may be multiple intermediate pairs with the same key and you can remember the current key. We are guaranteed to have all intermediate pairs with the same key end up on the same reducer. However, you can't assume that pairs with different keys will end up on the same reducer and therefore cannot remember any different keys that happened in the past. Doing it this way allows us to run map and reduce in parallel on many machines. Because Hadoop will eventually sort our intermediate data after mapping, it is ok to sort the intermediate data using the built-in sort command at the terminal.

Because this concept is important, I will expand upon it here a bit more. *You are not allowed to use the information from the first input data to affect the output for the second input data.*

Example:

Think back to the word count problem. Let's say the word "cat" appeared two times in a row. You'll recall our intermediate values were **(cat, 1)** and **(cat, 1)**. It would be *wrong* for the mapper to output **(cat, 2)** because *the second time we see cat, we would be remembering that we just saw cat and this is not stateless*. We are relying on the fact that cat was just seen in order to output **(cat, 2)**. This is not how MapReduce works.

Therefore, we must output **(cat, 1)** when we see the first instance of "cat" and **(cat, 1)** when we see the second instance of "cat". If that isn't clear, think of it this way: what happens if your input file is split right between the two cat s and 1 cat is sent to one processor and the other cat is sent to another processor? There is no communication between the processors, so the second one would have no idea what the first one is producing. Therefore, the mappers would each produce a **(cat, 1)** pair.

The bottom line is this: you should not rely on previously read in information to affect your current output. If you are unsure what you are doing is allowed, ask me and I'll tell you.

Project Tasks

You must write 1 mapper file and 1 reducer file to solve each of the following problems. Each mapper and reducer must be in its own file for each problem. You must name them **mapperX.py** and **reducerX.py** (where **X** is the problem number).

You should end up with **3 mapper files** and **3 reducer files** for this project. For all of these problems, you must read in from the command line. Refer to the **`sys.stdin.readline()`** function for reading in from the command line. The MapReduce example in the slides goes over a couple of examples. In order to output to the command line, use the **`print(...)`** function (again, see slides for examples).

You are only allowed to write 1 mapper and 1 reducer per problem.

I am assuming that many of you are familiar with Python from the course prerequisites. You should have done some basic input and output, and that is all you will be doing in this project. You will read in from the command line and output to the command line. *Do not read in from specific files or output to specific files.*

Typing in a ton of data at the command line is a bit tedious. Instead, you'll want to redirect your input from a file and then redirect your output to a file. We are using input and output redirection. All commands will work whether you are on a Windows machine or a Linux machine.

Example:

Let's say your mapper program for problem 1 is called **mapper1.py** and you want to read in from a file called **input.txt**. You need to output your intermediate data to some file. Let's output the intermediate (key, value) pairs to a file called **intermediate.txt**. To do this, you would do the following:

```
python mapper1.py < input.txt > intermediate.txt
```

The **<** operator is a redirect operator. It is telling your OS: I want to send everything from **input.txt** to this python program. Similarly, the **>** is redirecting your print statements to a file called **intermediate.txt** instead of printing it to the command line. Your intermediate (key, value) pairs are stored in a file called **intermediate.txt**.

Since Hadoop will eventually sort our data for us, we need to rely on the OS to do the sorting for us. In order to use this, you must have your key come first in the intermediate output. The **sort** function will sort the lines in alphabetical order. Therefore, *make sure each (key, value) pair is on its own line*. To sort your intermediate data, use the following command:

```
sort < intermediate.txt > intersorted.txt
```

Let's assume your reducer is called **reducer1.py** and you want to output your answer to **output.txt**. You would do the following:

```
python reducer1.py < intersorted.txt > output.txt
```

The reason we are splitting them this way is because eventually we will be using Hadoop's streaming feature. If you write your code in this fashion, you'll be able to use the code with Hadoop without having to change a thing. It is also a good way to reinforce that the mapping phase and reducer phase are completely separate phases.

Problem 1: Words with same number of vowels (10 points)

This problem is similar to the problem we worked on in lecture with a small twist. Instead of printing out how many times a word appears in the file, you want to print out how many words have the same number of vowels. For this problem, only the number of vowels matter. The actual vowel is not important. The output will be the number of vowels (let's call it **x**), followed by a colon, followed by the number of words that had **x** vowels. The output is sorted by the number of vowels.

Example:

Let's say the file contained the following words:

hello how are you doing today hello

The output would be:

1 : 1

2 : 4

3 : 2

Words are separated by spaces, tabs and new lines only. As an example, "half-time" would be considered 1 word. The format should be as you see above: the number of vowels, followed by a space, a colon and then a space, followed by the number of words that contained that many vowels.

Using the example above, since 'how' contains 1 vowel, the first output is

1 : 1 (i.e. exactly 1 vowel occurred in exactly 1 word).

'hello', 'are', 'doing' and 'hello' contained 2 vowels each. This is why we ended up with **2 : 4**. A word appearing more than once will be counted more than once.

Finally, 'today' and 'you' contained 3 vowels. For this problem, we will assume 'y' is a vowel. The final output line was **3 : 2**.

The output doesn't necessarily need to be sorted by the number of vowels.

Problem 2: Words with exact same vowels (15 points)

The second problem will be a slight modification of the first one. This problem will output the number of words that contained the exact same vowels.

Example:

'hello' and 'pole' both contain exactly 1 e and exactly 1 o. The order of the vowels does not matter.

Imagine the following example:

hello moose pole

We would end up with the following output:

eo : 2

eeo : 1

We will also assume 'y' is a vowel for this problem.

The format should be as seen above: the vowels of each line are in alphabetical order, followed by a space, a colon, a space, then followed by the number of words that contained exactly those vowels. The keys don't have to be in alphabetical order in the final output. For example, "eiu" could come before "aai".

Problem 3: Clothing coordination (25 points)

This is a very difficult problem that will require some thought. You will be working through a problem that recommender sites deal with. Imagine an online store that sells clothing. The website has a list of items that pair well together. Let's say shirt number 1 is usually bought with pants number 6 or shoes number 10. If there is only 1 item in the shopping cart, it is easy to recommend those pants or those shoes and many websites do this. However, let's say a user wants to buy shirt number 1 and shirt number 2. What are the items that pair well with *both* 1 and 2? This is the motivation behind this problem: a recommender system based on multiple input values.

The input will be constructed in the following fashion with the following list for 1 particular item *appearing on one line (disregard the line break in this document)*:

ClothingID : ClothingIDOne ClothingIDTwo ClothingIDThree... !
NotClothingOneID NotClothingTwoID NotClothingThreeID...

Each item list will be on its own line (see example below). All of the items that pair well with **ClothingID** are listed before the exclamation point. All of the items that don't pair well with ClothingID are listed after the exclamation point. All IDs will be integer values to keep things simple. If clothing item 1 pairs well with clothing item

2, then clothing item 2 also pairs well with clothing item 1. The goal is this: produce all combinations of clothing items and the clothing items that they all pair well with. The following pages walk you through examples and output.

As a starting point, consider the example input and starting output shown below. The next few pages go into more detail about how to interpret the output.

Example Input:

```
1 : 2 3 5 6 7 ! 4 8 9
2 : 1 3 4 5 8 9 ! 6 7
3 : 1 2 4 9 ! 5 6 7 8
4 : 2 3 7 8 9 ! 1 5 6
5 : 1 2 6 8 9 ! 3 4 7
6 : 1 5 7 8 ! 2 3 4 9
7 : 1 2 3 4 5 6 ! 8 9
8 : 2 4 5 6 9 ! 1 3 7
9 : 2 3 4 5 8 ! 1 6 7
```

Initial Lines of Output:

```
1 2 : 3 5
1 2 3 :
1 2 3 4 :
1 2 3 4 5 :
1 2 3 4 5 6 :
1 2 3 4 5 6 7 :
1 2 3 4 5 6 7 8 :
1 2 3 4 5 6 7 8 9 :
1 2 3 4 5 6 7 9 :
1 2 3 4 5 6 8 :
...
```

The first output line is saying that, if 1 and 2 are in the shopping cart, then a reasonable recommendation is to buy 3 and 5. Here's why:

Example Input:

1	:	2	3	5	6	7	!	4	8	9
2	:	1	3	4	5	8	9	!	6	7
3	:	1	2	4	9	!	5	6	7	8
4	:	2	3	7	8	9	!	1	5	6
5	:	1	2	6	8	9	!	3	7	4
6	:	1	5	7	8	!	2	3	4	9
7	:	1	2	3	4	5	6	!	8	9
8	:	2	4	5	6	9	!	1	3	7
9	:	2	3	4	5	8	!	6	7	1

...in the ...

Initial Lines of Output:

1 2 : 3 5

1 2 3 :

1 2 3 4 :

4

The first line of output states that items 1 and 2 both pair well with 3 and 5 because...

...in the input, we see 1 pairs with 3 and 5
well, and 2 also pairs with 3 and 5.

1 2 3 4 5 6 8 :

...

Now examine the second output line:

Example Input:

1	:	2	3	5	6	7	!	4	8	9
2	:	1	3	4	5	8	9	!	6	7
3	:	1	2	4	9	!	5	6	7	8
4	:	2	3	7	8	9	!	1	6	
5	:	1	2	6	8	9	!	3	4	
6	:	1	5	7	8	!	2	3	4	
7	:	1	2	3	4	5	6			
8	:	2	4	5	6	9				
9	:	2	3	4	5	8				

Initial Lines of Output:

1	2	:	3	5
1	2	3	:	

The second line shows that 1, 2, and 3 pair with nothing because...

...in the input:

- 1 pairs with 5, 6, 7
- 2 pairs with 4, 5, 8, 9
- 3 pairs with 4, 9

No single item is common to all three sets and therefore the set of clothes 1, 2, and 3 doesn't pair well with anything.

Notice that we don't have to specify items the shopper is going to buy together anyway:

Example Input:

1	:	2	3	5	6	7	!	4	8	9
2	:	1	3	4	5	8	9	!	6	7
3	:	1	2	4	9	!	5	6	7	8
4	:	2	3	5	6	7	!	1	5	6

Even though 1, 2, and 3 pair well together, we don't specify that in the second line of output because...

Initial Lines of Output:

1 2 : 3 5

1 2 3 :

1 2 3 4 :

1 2 3 4 5

2 3 4 5

2 3 4

2 3 4

...items on the right side of the : should never include items on the left since the user is already going to buy everything on the left. There is no need to recommend something the user is going to buy anyway.

Since the full output is quite long, I will not be posting it here. Below, I've posted the next several output lines where there is actually a pairing:

Example Input:

```
1 : 2 3 5 6 7 ! 4 8 9
2 : 1 3 4 5 8 9 ! 6 7
3 : 1 2 4 9 ! 5 6 7 8
4 : 2 3 7 8 9 ! 1 5 6
5 : 1 2 6 8 9 ! 3 4 7
6 : 1 5 7 8 ! 2 3 4 9
7 : 1 2 3 4 5 6 ! 8 9
8 : 2 4 5 6 9 ! 1 3 7
9 : 2 3 4 5 8 ! 1 6 7
```

Selected Output of Pairings:

```
1 2 4 : 3
1 2 4 7 : 3
1 2 4 7 9 : 3
1 2 4 9 : 3
1 2 6 : 5
1 2 6 7 : 5
1 2 6 7 8 : 5
1 2 6 7 8 9 : 5
1 2 6 7 9 : 5
1 2 6 8 : 5
1 2 6 8 9 : 5
1 2 6 9 : 5
1 2 7 : 3 5
```

The combinations do not necessarily have to be in sorted order as seen above (e.g. 1 2 4 coming before 1 2 4 7). However, *the values on either side of the colon should be in sorted order* (e.g. 1 2 5 6 is correct, 1 5 2 6 is not). Here is a complete example of input:

Complete Example

Input:

```
1 : 2 3 5 ! 4
2 : 1 3 4 5 !
3 : 1 2 4 ! 5
4 : 2 3 ! 1 5
5 : 1 2 ! 3 4
```

Output:

```
1 2 : 3 5
1 2 3 :
1 2 3 4 :
1 2 3 4 5 :
1 2 3 5 :
1 2 4 : 3
1 2 4 5 :
1 2 5 :
1 3 : 2
1 3 4 : 2
1 3 4 5 : 2
1 3 5 : 2
1 4 : 2 3
1 4 5 : 2
1 5 : 2
2 3 : 1 4
2 3 4 :
2 3 4 5 :
2 3 5 : 1
2 4 : 3
2 4 5 :
2 5 : 1
3 4 : 2
3 4 5 : 2
3 5 : 1 2
4 5 : 2
```

The left hand side should be a list of the clothing IDs separated by a space. This is followed by a colon. The right hand side contains a sorted list of clothing IDs that pair well with everything on the left hand side. These are also separated by a space.

Submitting Your Work

When you are finished, submit *a .zip file* to the **Project 3 dropbox** that includes:

- Your 3 mappers
- Your 3 reducers